

Functional Reqs:

- Chrome Extension that sends an additional API request to backend whenever "like" or "dislike" is clicked
- Should be toggled clicks with some way to "track" that a user has clicked like or dislike on a video already
- Don't want users to have to explicitly log in and authenticate with our service
- We want the like/dislike count to be very close in its estimation, but it doesn't need to be real time or accurate to the order of seconds/milliseconds
- We should always be available or have really good fault tolerance/recovery since estimating the dislike count if we are down will make things very difficult to maintain correctness
- Users should not be able to spam or click "like" / "dislike" more than once
- Latency at displaying the "like" / "dislike" count can be a bit slow since the component can render after the page/video loads

Functional Apps Cont:

- Cannot store sensitive user info like their email / phone number

Non-Objectives:

- Don't need to deal with storing the information that a user clicked "like" or "dislike" on youtube video
- Don't need to deal with user authentication

Parameters

Total Active YouTube Users Per Day = 122 million

$$= \frac{122 \text{ million}}{24 \cdot 60} = 84,722.22 \text{ users per min}$$

$$= \frac{84,722.22}{60} \approx 1,412.04 \text{ users per second}$$

Assuming each user clicks either "like" or "dislike" this means there is approximately 1,412.04 reps per second more likely something in on order of mag = 4 greater "very unlikely"

This can be better modeled with a Poisson Dist.

with parameter $\lambda = 1412$ meaning that in the span of one second \Rightarrow on interval on average we will see 1412 repects a second

It is a pretty big assumption to say that each user is going to like/dislike every video they watch, but I think it is a fairly accurate estimate and overestimating it will not lead to too many negative consequences

Parameters Cont.

Consistency - Doesn't need to be absolute consistency but it should be fairly strong

Availability - We want high availability since we are not the owners of youtube any amount of downtime will lead to inaccurate "like"/"dislike" counts

Partitioning - Dependent on number of users / and size of the data / approximate total storage

Latency - No hard requirements so I will set this to 5 seconds

Data Model

At the very least we need these models

Schema:

Video id (VARCHAR, 50)

Like Count (BIGINT) →

Dislike Count (BIGINT) →

Highest like count for a video is 50 million so this will be more than enough

If using a relational database:

video id = 50 bytes
like count = P bytes
Dislike count = P bytes

) 65 bytes per video entry

A stat said there are 800 million YouTube videos.
Assuming a higher upper bound of 1 billion
videos $65 \text{ bytes} \cdot 1 \text{ billion} \approx 65 \text{ GB of storage}$
(ignoring metadata/indexes)

Data Model Cont

If using a non-relational DB like Redis

Key	Value
video id + like	# of likes
video id + dislike	# of dislikes

Each of the key/values are variable size but the size can be modeled with probability. A stat says a good goal is a like 90% of 4 to be seen on a video. Can assume the dislike 9%. Let's assume avg youtube video gets 500,000 views.

$$0.9 \cdot 500,000 = 20000 = 5 \text{ digits}$$

$$0.01 \cdot 500,000 = 5000 = 4 \text{ digits}$$

$$\text{video id + like} = 56 \text{ bytes}$$

$$\text{video id + dislike} = 59 \text{ bytes}$$

$$\text{like count} \approx 5 \text{ bytes}$$

$$\text{dislike count} \approx 4 \text{ bytes}$$

$$124 \text{ bytes} \times 1 \text{ billion} = \\ 124 \text{ GB of storage}$$

Because you are able to save so much more bytes with relational, we will use that

Throughput

Earlier we estimated that on average we would see 1412 req/s. Assuming each req/res uses 8 bytes $\rightarrow 1412 \cdot 8 = 11296 \text{ Bytes/sec}$
 $= 11.296 \text{ KB/sec}$

The max theoretical throughput $\approx 2.376 \text{ Mbit/sec}$
 $= 297.886 \text{ KB/s}$

This means that theoretically one application server can handle the traffic
Let's assume there is 50% reads / 50% writes
on each request \Rightarrow 706 reads/s
706 writes/s

Online it says a single mySQL node handles about 200 writes/s and the theoretical ceiling is 1000 writes/s with SSDs. To make this easier on the load we can add a replica/batch up the write requests and add a queue to store the messages.

Throughput Cont

To use the read throughput, we can add a replication layer to the database for reads, only especially since consistency is not a huge issue

Fault Tolerance

Since we are aiming for high availability we always want a hot standby ready in the case that the current one goes down, maybe even 2

Now we want a policy on DNS to swap the record entries in the event something goes down

Every 5 minutes, create a snapshot backup of the data in the currently used DB node in case something fails.

Worst case loss $5.66 \cdot 200 = 11320$ writes, but we can queue to get over that

Security

Since it is a public facing API and we don't have authentication, we need some form to prevent spamming

1. To prevent DDoS attacks, should put a WAF in front of the application server that limits people out it may send more than 5 req/s
2. Have a cache that stores 10 fb worth of user cookies to their state

Only happens on writes so this has 76 req/s which a single redis node can easily handle

If a user recently wrote "like" a video, the state is stored in cache that the App servers check. If the user is in the "like" state, it prevents "liking" again. Same for "dislike"

Overall Design

Based on the previous pages, our final design

looks like this:

